

Case_Study_4_Code

Brian Gaither, Sean Mcwhirter, Andrew Mejia, Sabrina Purvis

2021-01-31 20:27:13

Contents

Functions	4
The gen_Table function will parse through the table of 20 records and 15 observations	5
and parse the table into its own data frame from xml2 read)html function and then use	5
the pipe operator to use rvest nodes to find the table structure	5
then the fucntion will insert the metadata for the query parameters of year, division, section, page, source link and sex	5
This function will use all available cores on machine	6
It will process the years in parrell.	6
This code has been adapted from	6
https://github.com/ngupta23/ds7333_qtw/blob/master/case_study_2/submission_Kannan_Moro_Gupta/code/CS2_ETL.Rmd	6
https://cran.r-project.org/web/packages/doSNOW/doSNOW.pdf	6
https://stackoverflow.com/questions/36794063/r-foreach-from-single-machine-to-cluster	6
https://cran.r-project.org/web/packages/progress/progress.pdf	6
https://www.r-bloggers.com/2013/08/the-wonders-of-foreach/	6
The purpose of this function is to transform the raw tables from the scrape	7
https://stackoverflow.com/questions/50040968/convert-a-duration-hms-to-seconds	7
https://stackoverflow.com/questions/10835908/is-there-a-way-to-convert-mmss-00-to-seconds-00	7
https://stackoverflow.com/questions/24173194/remove-parentheses-and-text-within-from-strings-in-r	7

Perform Scrape	8
Caution takes a long time without hexacore machine	8
Load tables from file	9
Preview of raw table scrapes	9
Perform the table transformation and remove metadata columns and other columns	10
Preview of the transformed table	10
Preview of NA columns	11
Create a DF containing the NA counts of each feature	12
Plots of prominate NA columns	12
We will remove NAs and only focus on completed records and rewmove NAs for further processing and reindex the table.	14
Final dataframe metadata	14
Generate plots of quick EDA for ages accross years	15
Conclusion / Recommendations	36

```
library(tidyverse)
```

```
## -- Attaching packages ----- tidyverse 1.3.0 --
```

```
## v ggplot2 3.3.3      v purrr  0.3.4
## v tibble  3.0.4      v dplyr  1.0.2
## v tidyr   1.1.2      v stringr 1.4.0
## v readr   1.4.0      v forcats 0.5.0
```

```
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()    masks stats::lag()
```

```
library(hms)
library(rvest)
```

```
## Loading required package: xml2
```

```
##
```

```
## Attaching package: 'rvest'
```

```
## The following object is masked from 'package:purrr':
```

```
##
```

```
##      pluck
```

```
## The following object is masked from 'package:readr':
```

```
##
```

```
##      guess_encoding
```

```
library(lubridate)
```

```
##  
## Attaching package: 'lubridate'  
  
## The following object is masked from 'package:hms':  
##  
##     hms  
  
## The following objects are masked from 'package:base':  
##  
##     date, intersect, setdiff, union
```

```
library(foreach)
```

```
##  
## Attaching package: 'foreach'  
  
## The following objects are masked from 'package:purrr':  
##  
##     accumulate, when
```

```
library(stringr)  
library(iterators)  
library(progress)  
library(doParallel)
```

```
## Loading required package: parallel
```

```
library(doSNOW)
```

```
## Loading required package: snow  
  
##  
## Attaching package: 'snow'  
  
## The following objects are masked from 'package:parallel':  
##  
##     clusterApply, clusterApplyLB, clusterCall, clusterEvalQ,  
##     clusterExport, clusterMap, clusterSplit, makeCluster, parApply,  
##     parCapply, parLapply, parRapply, parSapply, splitIndices,  
##     stopCluster
```

```
library(dplyr)  
library(states)
```

```
##  
## Attaching package: 'states'
```

```
## The following object is masked from 'package:readr':
##
##   parse_date

library(ggplot2)
library(ggthemes)
library(SiZer)
library(plotly)

##
## Attaching package: 'plotly'

## The following object is masked from 'package:ggplot2':
##
##   last_plot

## The following object is masked from 'package:stats':
##
##   filter

## The following object is masked from 'package:graphics':
##
##   layout
```

```
years = c(1999:2012)
division = 'Overall+Women'
section = '10M'
sex = 'W'
```

Functions #The gen_Link function will generates the link with the query parameters for the searchable database

```
gen_Link = function(year,division,section,page=1,sex){

  paste0( 'http://www.cballtimeresults.org/performances'
    , '?division=',division,'&page=',page,
    '&section=',section, '&sex=',sex,
    '&utf8=%E2%9C%93', '&year=',year)
    #, '?utf8=%E2%9C%93&section=',section
    #, '&year=',year, '&division=',division, '&page=', page)

}
```

The `gen_Table` function will parse through the table of 20 records and 15 observations

and parse the table into its own data frame from `xml2::read_html` function and then use

the pipe operator to use `rvest` nodes to find the table structure

then the function will insert the metadata for the query parameters of year, division, section, page, source link and sex

```
gen_Table = function(year,division,section,page, sex){  
  
  #use gen_link function to get link to page  
  genlink=gen_Link(year,division,section,page=page, sex=sex)  
  
  #read the page, and grab to 'table' tag  
  single_table = xml2::read_html(genlink) %>%  
    rvest::html_nodes("table") %>%  
    rvest::html_table(fill=TRUE)  
  
  #get the table and add metadata for the query parameters  
  table_out = single_table[[1]] %>%  
    mutate(year=year, divisionTitle=division, section=section, page=page, source=genlink, sex = sex)  
}
```

This function will use all available cores on machine

It will process the years in parallel.

This code has been adapted from

https://github.com/ngupta23/ds7333_qtw/blob/master/case_study_2/submission_Kannan_Moro_Gupta/code/CS2_ETL.Rmd

<https://cran.r-project.org/web/packages/doSNOW/doSNOW.pdf>

<https://stackoverflow.com/questions/36794063/r-foreach-from-single-machine-to-cluster>

<https://cran.r-project.org/web/packages/progress/progress.pdf>

<https://www.r-bloggers.com/2013/08/the-wonders-of-foreach/>

```
scrapeTables = function(years,division,section, sex, max_itr = 500){

library(progress)
library(doParallel)
library(doSNOW)

#Initialize Parallel Process to detect number of cores
#https://cran.r-project.org/web/packages/doSNOW/doSNOW.pdf
#https://stackoverflow.com/questions/36794063/r-foreach-from-single-machine-to-cluster

#Generate and register initial clusters based on cores, otherwise, this is a long process
cl = makeCluster(detectCores())
doSNOW::registerDoSNOW(cl)

#Generate progress bar for the parallel loop based on number of years
#https://cran.r-project.org/web/packages/progress/progress.pdf
progBar = progress::progress_bar$new(total = length(years),format='[:bar] :percent :eta')
progress = function(n) progBar$tick()

#Initialize a parallel loop per each year
#Initialize tableRaw as empty to loop to be populated as a table dataframe from gen_Table function
tableRaw=NULL

#tableRaw will now use for each using years as the iterator to use .combine to rbind
#.export will do the gen_Table and gen_Link functions simultaneously and
```

```

#options.snow will show the progress bar
# the %dopar% will process all the years simultaneously.
#https://www.r-bloggers.com/2013/08/the-wonders-of-foreach/
tableRaw = foreach(y=years
                    ,.combine=rbind,.export=c('gen_Table','gen_Link')
                    ,.options.snow = list(progress=progress)) %dopar%
{
  library(foreach)
  library(dplyr)
  #initialize isCompleted variable as FALSE for bool conditions to see if loop has been completed
  isCompleted=FALSE

  #Initiate loop since most pages are 487, we will only loop for the iterations for max_itr
  tableRaw=foreach(p=c(1:max_itr),.combine=rbind) %do%
  if(!isCompleted) {
    message('getting year:',y, ' page:',p,appendLF = F)
    #get the table of the current page
    table = gen_Table(year=y
                      ,division=division
                      ,section=section
                      ,page=p
                      ,sex=sex)
    message(' rows:',nrow(table))
    isCompleted = nrow(table)==0 #if there is record, we are at the last page, no need to read
    return(table)
  }
  return(tableRaw)
}

#Deactivate the cluster of cores
stopCluster(cl)
#save the raw data to rda format for later processing based on gender
saveRDS(tableRaw,file=paste0('CB',sex,'tableRaw.rds'))

return(tableRaw)
}

```

The purpose of this function is to transform the raw tables from the scrape

<https://stackoverflow.com/questions/50040968/convert-a-duration-hms-to-seconds>

<https://stackoverflow.com/questions/10835908/is-there-a-way-to-convert-mmss-00-to-seconds-00>

<https://stackoverflow.com/questions/24173194/remove-parentheses-and-text-within-from-strings-in-r>

```

tableTransform =function(data_df, cols_to_remove=NULL){
  dataDF = data_df %>%
    #Seperate Home town into seperate columns
    separate(col = 'Hometown', c('HomeTown', 'HomeState'), sep = ',', extra = 'merge', remove = TRUE, f
    #Seperate PiS/TiS into seperate columns

```

```

separate(col='PiS/TiS',c('PiS','TiS'),sep='\\/'
,extra='drop',remove=TRUE) %>%
#Seperate PiD/TiD into seperate columns
separate(col='PiD/TiD',c('PiD','TiD'),sep='\\/'
,extra='drop',remove=TRUE) %>%
#Trim the casted upper HomeTown strings of whitespace
mutate(HomeTown = toupper(trimws(HomeTown))
, HomeState = toupper(trimws(HomeState))
, HomeTown = ifelse(HomeTown %in% c('NR', '', NULL), NA, toupper(trimws(HomeTown)))
, HomeState = ifelse(HomeState %in% c('NR', '', NULL), NA, toupper(trimws(HomeState)))
#Check if HomeState is in state.abb or DC and return USA else return the HomeTown as the Country
, HomeCountry = ifelse(HomeState %in% c(state.abb, "DC"), "USA", HomeTown)
#Remove White Space
, PiS = ifelse(trimws(PiS) %in% c('NR', '', NULL), NA, trimws(PiS))
, TiS = ifelse(trimws(TiS) %in% c('NR', '', NULL), NA, trimws(TiS))
, PiD = ifelse(trimws(PiD) %in% c('NR', '', NULL), NA, trimws(PiD))
, TiD = ifelse(trimws(TiD) %in% c('NR', '', NULL), NA, trimws(TiD))
, Division = ifelse(trimws(Division) %in% c('NR', '', NULL), NA, trimws(Division))
# Normalize Time to seconds and minutes
, RawTime = strptime(Time, format='%H:%M:%S')
, RawTime_S = RawTime$hour * 3600 + RawTime$min * 60 + RawTime$sec
, RawTime_M = as.numeric(RawTime_S)/60
, RawPace = strptime(Pace, format = "%M:%OS")
, RawPace_S = RawPace$min * 60 + RawPace$sec
#Normalize Age where 'NR' as NA
, Age = ifelse(Age %in% c("NR"), NA, Age)
#Cast Variables as appropriate dtypes
, Age = as.numeric(Age)
, RawTime_S = as.numeric(RawTime_S)
, RawPace_S = as.numeric(RawPace_S)
, RawPace_M = as.numeric(RawPace_S)/60
, year = as.factor(year)
#Remove the (<Sex>) from the names
#, Name = str_replace(Name, " \\s*\\([~\\|\\|]+\\|\\)", '')
)
#Remove columns we do not want
dataDF = dataDF %>% select (-all_of(c(cols_to_remove)))

return(dataDF %>% select(c( Age, year, HomeTown, HomeState, HomeCountry, RawTime_S, RawTime_M, RawPace_M, RawPace_S)))
}

```

Perform Scrape

Caution takes a long time without hexacore machine

```

#Women_table = scrapeTables(years=years,division = division,section=section,sex = sex, max_itr = 500)

#men_table = scrapeTables(years=years,division = 'Overall+Men',section=section,sex = "M", max_itr = 500)

```


Load tables from file

```
mens_table <- readRDS("/media/andrew/Seagate Backup Plus Drive/Documents/School/HomeWork/QTW/DS7333/CASI
womens_table <- readRDS("/media/andrew/Seagate Backup Plus Drive/Documents/School/HomeWork/QTW/DS7333/C
```

Preview of raw table scrapes

```
head(womens_table, n = 10)
```

##	Race	Name	Age	Time	Pace	PiS/TiS	Division	PiD/TiD
## 1	1999 10M	Jane Omoro (W)	26	0:53:37	5:22	1/2358	W2529	1/559
## 2	1999 10M	Jane Ngotho (W)	29	0:53:38	5:22	2/2358	W2529	2/559
## 3	1999 10M	Lidiya Grigoryeva (W)	NR	0:53:40	5:22	3/2358	NR	NR
## 4	1999 10M	Eunice Sagero (W)	20	0:53:55	5:24	4/2358	W2024	1/196
## 5	1999 10M	Alla Zhilyayeva (W)	29	0:54:08	5:25	5/2358	W2529	3/559
## 6	1999 10M	Teresa Wanjiku (W)	24	0:54:10	5:25	6/2358	W2024	2/196
## 7	1999 10M	Elana Viazova (W)	38	0:54:29	5:27	7/2358	W3539	1/387
## 8	1999 10M	Gladys Asiba (W)	NR	0:54:50	5:29	8/2358	NR	NR
## 9	1999 10M	Nnenna Lynch (W)	27	0:55:39	5:34	9/2358	W2529	4/559
## 10	1999 10M	Margaret Kagiri (W)	30	0:55:43	5:34	10/2358	W3034	1/529

##	Hometown	year	division	Title	section	page
## 1	Kenya	1999	Overall+Women	10M	1	
## 2	Kenya	1999	Overall+Women	10M	1	
## 3	Russia	1999	Overall+Women	10M	1	
## 4	Kenya	1999	Overall+Women	10M	1	
## 5	Russia	1999	Overall+Women	10M	1	
## 6	Kenya	1999	Overall+Women	10M	1	
## 7	Ukraine	1999	Overall+Women	10M	1	
## 8	Kenya	1999	Overall+Women	10M	1	
## 9	Concord, MA	1999	Overall+Women	10M	1	
## 10	Kenya	1999	Overall+Women	10M	1	

## 1	http://www.cballtimeresults.org/performances?division=Overall+Women&page=1&section=10M&sex=W&utf8=
## 2	http://www.cballtimeresults.org/performances?division=Overall+Women&page=1&section=10M&sex=W&utf8=
## 3	http://www.cballtimeresults.org/performances?division=Overall+Women&page=1&section=10M&sex=W&utf8=
## 4	http://www.cballtimeresults.org/performances?division=Overall+Women&page=1&section=10M&sex=W&utf8=
## 5	http://www.cballtimeresults.org/performances?division=Overall+Women&page=1&section=10M&sex=W&utf8=
## 6	http://www.cballtimeresults.org/performances?division=Overall+Women&page=1&section=10M&sex=W&utf8=
## 7	http://www.cballtimeresults.org/performances?division=Overall+Women&page=1&section=10M&sex=W&utf8=
## 8	http://www.cballtimeresults.org/performances?division=Overall+Women&page=1&section=10M&sex=W&utf8=
## 9	http://www.cballtimeresults.org/performances?division=Overall+Women&page=1&section=10M&sex=W&utf8=
## 10	http://www.cballtimeresults.org/performances?division=Overall+Women&page=1&section=10M&sex=W&utf8=

##	sex
## 1	W
## 2	W
## 3	W
## 4	W
## 5	W
## 6	W

```
## 7    W
## 8    W
## 9    W
## 10   W
```

Perform the table transformation and remove metadata columns and other columns

```
cols_for_remove = c("divisionTitle", "source", "Pace", "Time", "RawTime", "RawPace", "page", "Race", "s")
mens_table_T = tableTransform(mens_table, cols_to_remove = all_of(cols_for_remove))
```

```
## Warning: Expected 2 pieces. Missing pieces filled with 'NA' in 3 rows [42230,
## 69995, 69996].
```

```
## Warning: Expected 2 pieces. Missing pieces filled with 'NA' in 31 rows [1207,
## 2797, 7583, 9206, 9255, 10991, 12255, 13046, 14812, 15842, 21756, 21819, 22093,
## 22709, 23342, 23680, 25014, 25424, 25443, 25750, ...].
```

```
womens_table_T = tableTransform(womens_table, cols_to_remove = all_of(cols_for_remove))
```

```
## Warning: Expected 2 pieces. Missing pieces filled with 'NA' in 20 rows [3, 8,
## 17, 2176, 7135, 7766, 8777, 9680, 10831, 18391, 18399, 18981, 19694, 20735,
## 21480, 22189, 28388, 29223, 38455, 47506].
```

Preview of the transformed table

```
head(womens_table_T, n = 10)
```

```
##      Age year HomeTown HomeState HomeCountry RawTime_S RawTime_M RawPace_S
## 1    26 1999    KENYA      <NA>      KENYA      3217    53.61667      322
## 2    29 1999    KENYA      <NA>      KENYA      3218    53.63333      322
## 3    NA 1999    RUSSIA      <NA>      RUSSIA      3220    53.66667      322
## 4    20 1999    KENYA      <NA>      KENYA      3235    53.91667      324
## 5    29 1999    RUSSIA      <NA>      RUSSIA      3248    54.13333      325
## 6    24 1999    KENYA      <NA>      KENYA      3250    54.16667      325
## 7    38 1999  UKRAINE      <NA>      UKRAINE      3269    54.48333      327
## 8    NA 1999    KENYA      <NA>      KENYA      3290    54.83333      329
## 9    27 1999  CONCORD      MA        USA        3339    55.65000      334
## 10   30 1999    KENYA      <NA>      KENYA      3343    55.71667      334
##      RawPace_M sex section PiS  TiS  PiD  TiD
## 1    5.366667   W    10M    1 2358    1  559
## 2    5.366667   W    10M    2 2358    2  559
## 3    5.366667   W    10M    3 2358 <NA> <NA>
## 4    5.400000   W    10M    4 2358    1  196
## 5    5.416667   W    10M    5 2358    3  559
## 6    5.416667   W    10M    6 2358    2  196
## 7    5.450000   W    10M    7 2358    1  387
## 8    5.483333   W    10M    8 2358 <NA> <NA>
## 9    5.566667   W    10M    9 2358    4  559
## 10   5.566667   W    10M   10 2358    1  529
```

```
dim(womens_table_T)
```

```
## [1] 75866    15
```

```
summary(womens_table_T)
```

```
##      Age      year      HomeTown      HomeState
## Min.   : 7.00   2012   : 9727   Length:75866   Length:75866
## 1st Qu.:27.00   2011   : 9030   Class :character   Class :character
## Median :32.00   2010   : 8853   Mode  :character   Mode  :character
## Mean   :33.85   2009   : 8323
## 3rd Qu.:39.00   2008   : 6395
## Max.   :87.00   2007   : 5532
## NA's   :20      (Other):28006
## HomeCountry      RawTime_S      RawTime_M      RawPace_S
## Length:75866     Min.   : 3104   Min.   : 51.73   Min.   : 310.0
## Class :character  1st Qu.: 5319   1st Qu.: 88.65   1st Qu.: 532.0
## Mode  :character  Median : 5849   Median : 97.48   Median : 585.0
##                      Mean   : 5893   Mean   : 98.22   Mean   : 589.4
##                      3rd Qu.: 6418   3rd Qu.:106.97   3rd Qu.: 642.0
##                      Max.   :10651   Max.   :177.52   Max.   :1065.0
##
## RawPace_M      sex      section      PiS
## Min.   : 5.167   Length:75866   Length:75866   Length:75866
## 1st Qu.: 8.867   Class :character   Class :character   Class :character
## Median : 9.750   Mode  :character   Mode  :character   Mode  :character
## Mean   : 9.823
## 3rd Qu.:10.700
## Max.   :17.750
##
##      TiS      PiD      TiD
## Length:75866   Length:75866   Length:75866
## Class :character   Class :character   Class :character
## Mode  :character   Mode  :character   Mode  :character
##
##
##
##
```

Preview of NA columns

```
columns = c("Age", "year", "HomeTown", "HomeState", "HomeCountry", "RawTime_S", "RawTime_M", "RawPace_S")
womens_table_T_na = womens_table_T %>% filter_at(vars(all_of(columns)),any_vars(is.na(.)))
head(womens_table_T_na, n = 10)
```

```
##      Age year HomeTown HomeState HomeCountry RawTime_S RawTime_M RawPace_S
## 1    26 1999    KENYA    <NA>    KENYA      3217    53.61667      322
## 2    29 1999    KENYA    <NA>    KENYA      3218    53.63333      322
```

```
## 3 NA 1999 RUSSIA <NA> RUSSIA 3220 53.66667 322
## 4 20 1999 KENYA <NA> KENYA 3235 53.91667 324
## 5 29 1999 RUSSIA <NA> RUSSIA 3248 54.13333 325
## 6 24 1999 KENYA <NA> KENYA 3250 54.16667 325
## 7 38 1999 UKRAINE <NA> UKRAINE 3269 54.48333 327
## 8 NA 1999 KENYA <NA> KENYA 3290 54.83333 329
## 9 30 1999 KENYA <NA> KENYA 3343 55.71667 334
## 10 NA 1999 LANCASTER PA USA 3576 59.60000 358
## RawPace_M sex section PiS TiS PiD TiD
## 1 5.366667 W 10M 1 2358 1 559
## 2 5.366667 W 10M 2 2358 2 559
## 3 5.366667 W 10M 3 2358 <NA> <NA>
## 4 5.400000 W 10M 4 2358 1 196
## 5 5.416667 W 10M 5 2358 3 559
## 6 5.416667 W 10M 6 2358 2 196
## 7 5.450000 W 10M 7 2358 1 387
## 8 5.483333 W 10M 8 2358 <NA> <NA>
## 9 5.566667 W 10M 10 2358 1 529
## 10 5.966667 W 10M 17 2358 <NA> <NA>
```

```
dim(womens_table_T_na)
```

```
## [1] 262 15
```


Create a DF containing the NA counts of each feature

```
na_df = data.frame(rowsum+(is.na(womens_table_T)), womens_table_T$year))
na_df = cbind(year_idx = rownames(na_df), na_df)
```


Plots of prominate NA columns

```
p1_na = ggplot(na_df, aes(x=year_idx, y=HomeTown)) + geom_bar(stat = "identity", fill = "grey") + labs(
p1_na
```

Counts of missing HomeTown Records by Year



```
p2_na = ggplot(na_df, aes(x=year_idx, y=HomeState)) + geom_bar(stat = "identity", fill = "grey") + lab  
p2_na
```



We will remove NAs and only focus on completed records and rewmove NAs for further processing and reindex the table.

```
womens_table_T = womens_table_T[complete.cases(womens_table_T), ]
row.names(womens_table_T) = NULL
```

Final dataframe metadata

```
head(womens_table_T, n = 10)
```

##	Age	year	HomeTown	HomeState	HomeCountry	RawTime_S	RawTime_M	RawPace_S
## 1	27	1999	CONCORD	MA	USA	3339	55.65000	334
## 2	30	1999	EUGENE	OR	USA	3373	56.21667	337
## 3	37	1999	BLOOMINGTON	MN	USA	3443	57.38333	344
## 4	39	1999	ALBUQUERQUE	NM	USA	3444	57.40000	344
## 5	32	1999	CHAPEL HILL	NC	USA	3471	57.85000	347
## 6	30	1999	WASHINGTON	DC	USA	3485	58.08333	349
## 7	31	1999	COLUMBIA	MD	USA	3516	58.60000	352
## 8	25	1999	ALEXANDRIA	VA	USA	3582	59.70000	358
## 9	38	1999	SILVER SPRING	MD	USA	3588	59.80000	359
## 10	31	1999	ROCKVILLE	MD	USA	3630	60.50000	363

```
##      RawPace_M sex section PiS  TiS PiD TiD
## 1    5.566667  W     10M   9 2358  4 559
## 2    5.616667  W     10M  11 2358  2 529
## 3    5.733333  W     10M  12 2358  2 387
## 4    5.733333  W     10M  13 2358  3 387
## 5    5.783333  W     10M  14 2358  3 529
## 6    5.816667  W     10M  15 2358  4 529
## 7    5.866667  W     10M  16 2358  5 529
## 8    5.966667  W     10M  18 2358  5 559
## 9    5.983333  W     10M  19 2358  4 387
## 10   6.050000  W     10M  20 2358  6 529
```


```
Descriptions = c('Participant Age', 'Year of race', 'Participant Home Town', 'Participant Home State',
womensTableInfo = data.frame(Feature = names(womens_table_T), Description = Descriptions, Type = sapply(
womensTableInfo
```

##	Feature	Description	Type
## 1	Age	Participant Age	double
## 2	year	Year of race	integer
## 3	HomeTown	Participant Home Town	character
## 4	HomeState	Participant Home State	character
## 5	HomeCountry	Participant Home Country	character
## 6	RawTime_S	Participant's Time in Seconds	double
## 7	RawTime_M	Participant's Time in Minutes	double
## 8	RawPace_S	Participants Mile Pace in Seconds	double
## 9	RawPace_M	Participants Mile Pace in Minutes	double
## 10	sex	Gender	character
## 11	section	Participant Race	character
## 12	PiS	Position in Sex	character
## 13	TiS	Total in Sex	character
## 14	PiD	Position in Division	character
## 15	TiD	Total in Division	character

Generate plots of quick EDA for ages accross years


```
plot_data_age = ggplot(womens_table_T, aes(x = year, y = Age)) + geom_boxplot() + labs(title = "Distrib
plot_data_age
```

Distribution of Women Participants Age by Year



```
plot_data_time = ggplot(womens_table_T, aes(x = year, y = RawTime_M)) + geom_boxplot() + labs(title = "Distribution of Women Participants Age by Year")
plot_data_time
```



Distribution of Women Raw Time by Year



```
plot_data_Pace = ggplot(womens_table_T, aes(x = year, y = RawPace_M)) + geom_boxplot() + labs(title = "Distribution of Women Raw Time by Year")
```

```
plot_data_Pace
```

Distribution of Women Raw Pace by Year



BRIAN GAITHER #####
#####

Creating an age bin column for analysis

```
womens_table_T$AgeBin = cut(womens_table_T$Age, breaks=c(0,5,15,25,35,45,55,65,75,85,95),labels=c("1-5"
```

Creating a column to bin the pace in minutes for later analysis

```
womens_table_T$PaceBin = cut(womens_table_T$RawPace_M, breaks=c(0,5,5.5,6,6.5,7,7.5,8,8.5,9,9.5,10,10.5,  
labels=c("1-5", "5.1-5.5", "5.5-6", "6-6.5", "6.5-7", "7-7.5", "7.5-8", "8-8.5", "8.5-9", "9-9.5"
```

write out the dataframe for later retrieval


```
#write.csv(womens_table_T,"C:/Users/blgai/OneDrive/Documents/School/SMU/Spring 2021/Quantifying the Wo
```

read the data into dataframe

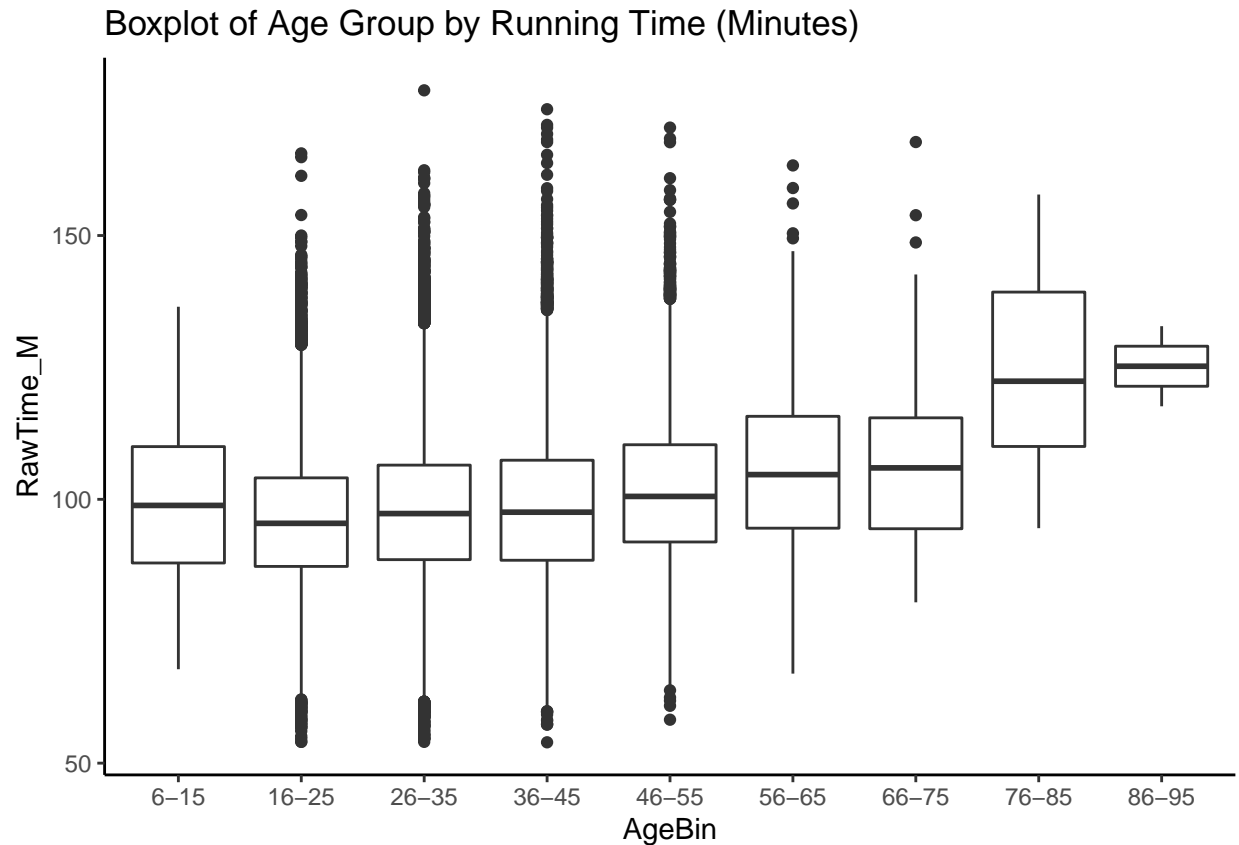
```
#womens_table_T = read.csv("C:/Users/blgai/OneDrive/Documents/School/SMU/Spring 2021/Quantifying the Wo
```

```
ggplot(womens_table_T, aes(x = Age, y=RawTime_M)) +  
  geom_hex() +  
  theme_bw() +  
  labs(title="Scatter Plot of Female Runners: Age vs Time in Minutes")
```

Scatter Plot of Female Runners: Age vs Time in Minutes



```
f = ggplot(womens_table_T)
f + geom_boxplot(mapping = aes(x=AgeBin, y=RawTime_M)) + theme_classic() + labs(title="Boxplot of Age G
```



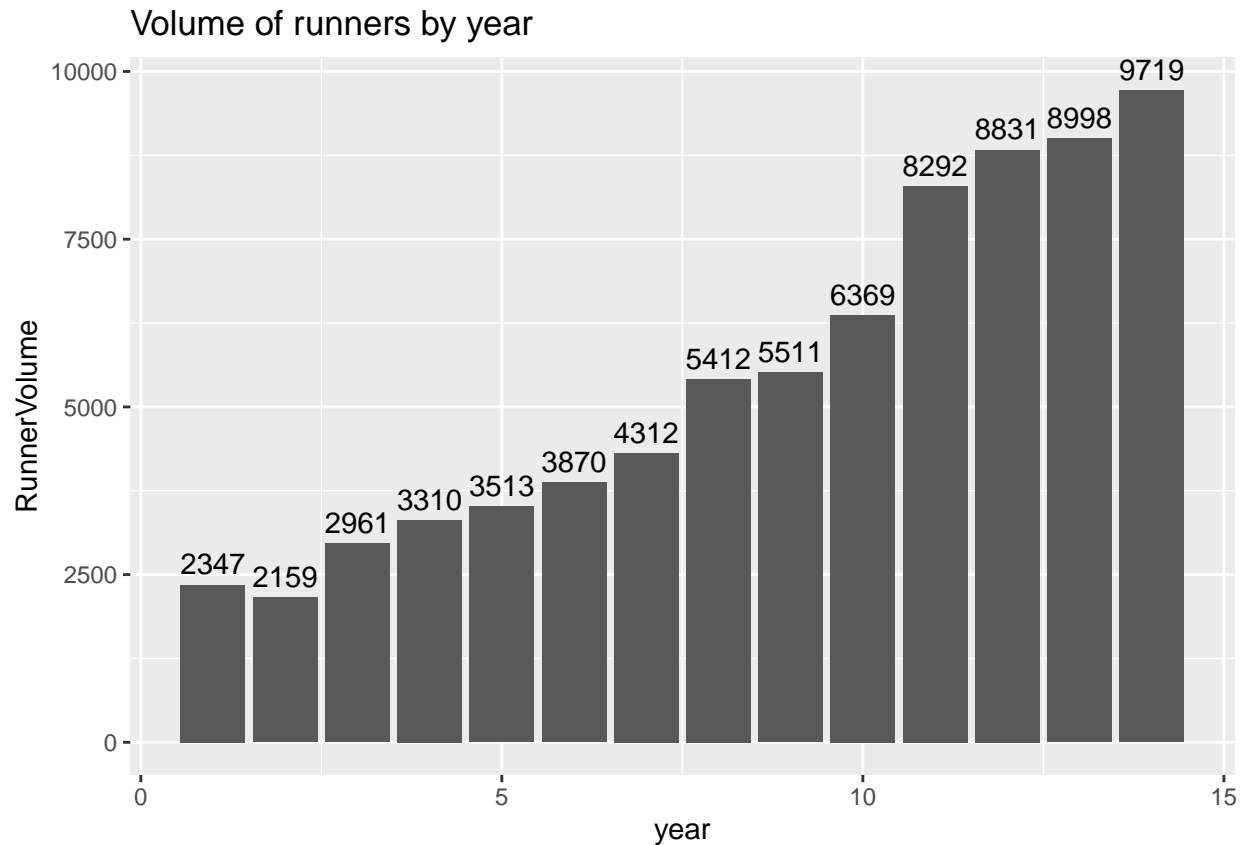
Let's examine overall runner volume

```
dfRunVol = womens_table_T %>% group_by(year) %>% tally()
dfRunVol$year = as.integer(dfRunVol$year)
names(dfRunVol)[2] <- "RunnerVolume"
head(dfRunVol)
```

```
## # A tibble: 6 x 2
##   year RunnerVolume
##   <int>      <int>
## 1     1        2347
## 2     2        2159
## 3     3        2961
## 4     4        3310
## 5     5        3513
## 6     6        3870
```

Plotting out volume of runners by year

```
ggplot(dfRunVol, aes(x=year, y=RunnerVolume)) +
  geom_col() +
  labs(title="Volume of runners by year") +
  geom_text(aes(label = RunnerVolume), vjust = -0.5)
```



let's see how well we can predict the volume of racers in the 15th year

```
lmRunVol = lm(RunnerVolume ~ year, data = dfRunVol)
lmRunVol$coefficients
```

```
## (Intercept)      year
##    800.1978    613.3451
```

Below we see we have a statistically significant slope and intercept

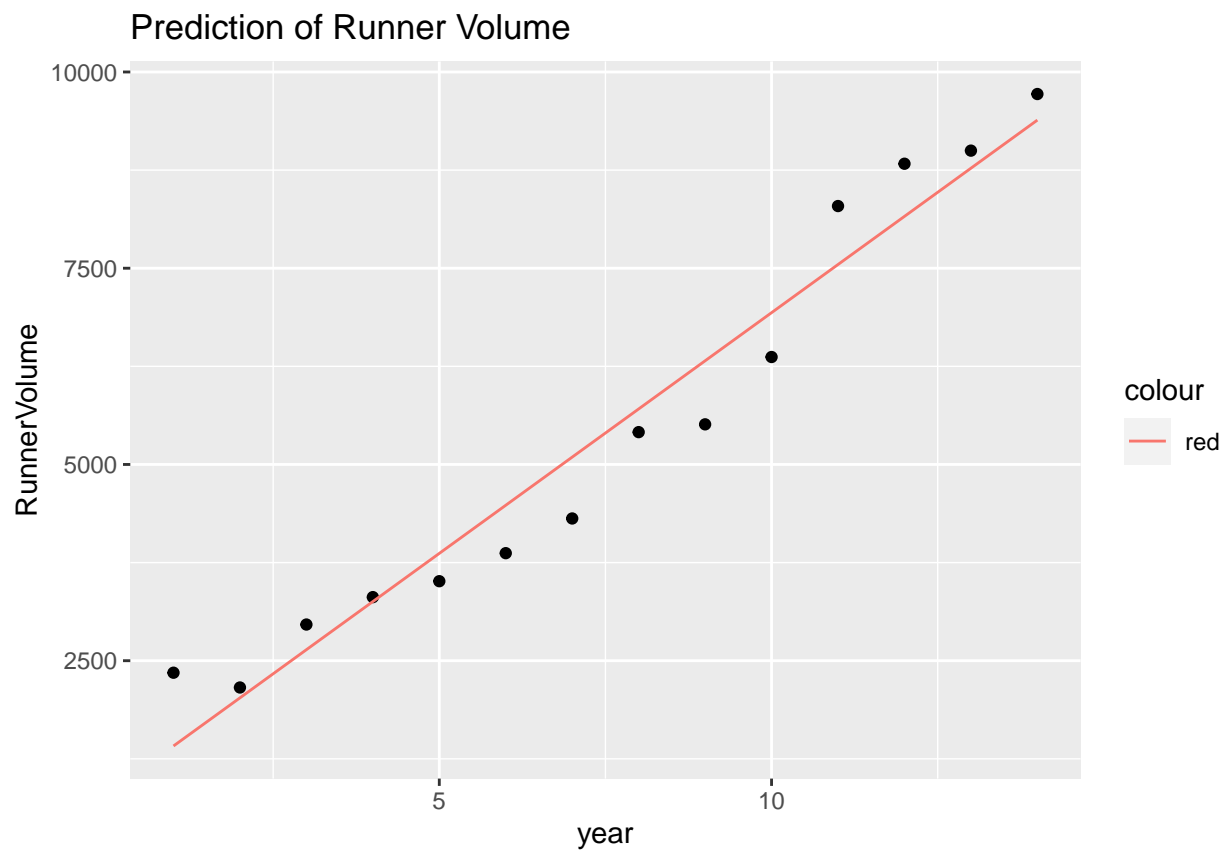
```
summary(lmRunVol)
```

```
##
## Call:
## lm(formula = RunnerVolume ~ year, data = dfRunVol)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -809.30 -511.97   94.27  329.17  933.46
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   800.20     339.09    2.36  0.0361 *
## year          613.35     39.82   15.40 2.87e-09 ***
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 600.7 on 12 degrees of freedom
## Multiple R-squared:  0.9518, Adjusted R-squared:  0.9478
## F-statistic: 237.2 on 1 and 12 DF,  p-value: 2.872e-09
```

See how well our line fits

```
preds = predict(lmRunVol)
ggplot(dfRunVol, aes(x=year, y=RunnerVolume)) + geom_point() + geom_line(dfRunVol, mapping = aes(x=year
```



How many runners are predicted to attend in 2013 (year 15)

```
y_hat = 800.5 + (15*613.4)
print(paste0("Predicted volume of runners in 2013 is: ", y_hat))
```

```
## [1] "Predicted volume of runners in 2013 is: 10001.5"
```

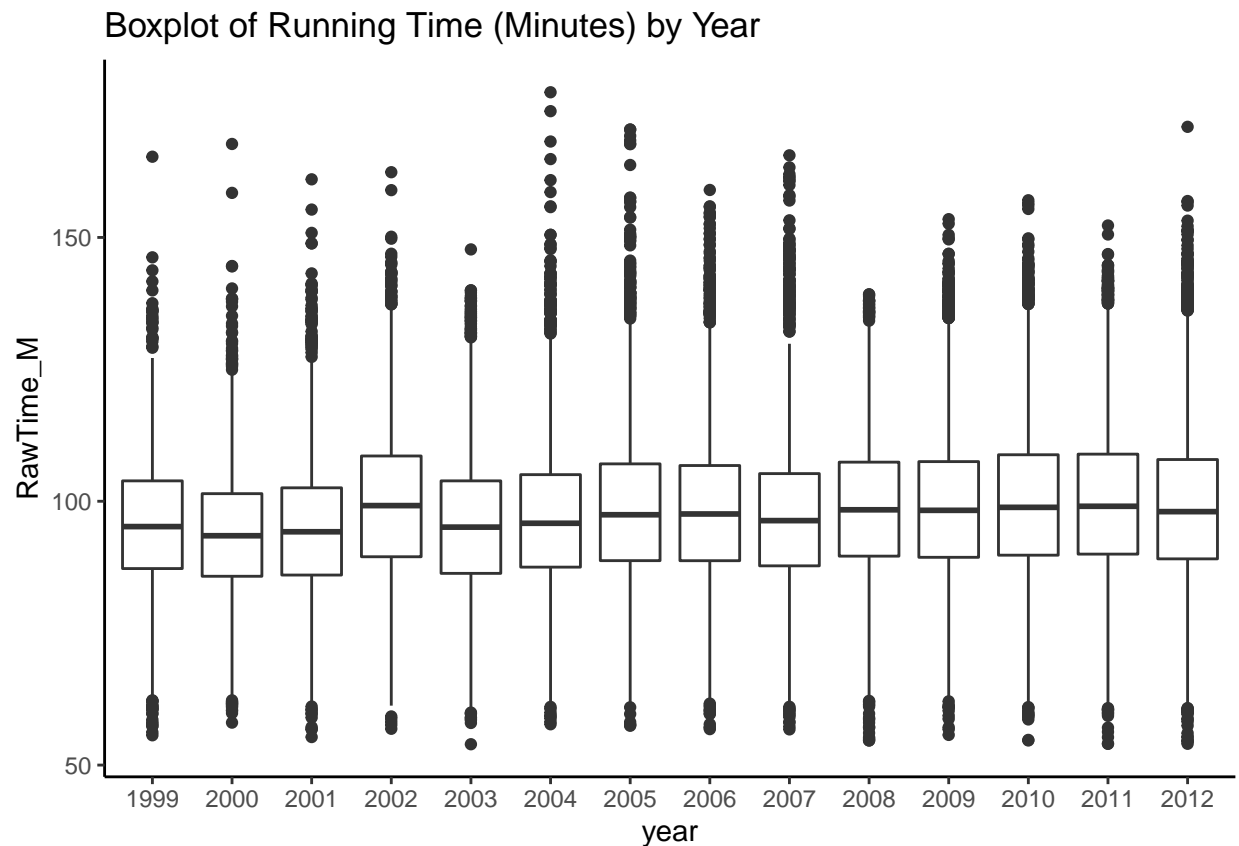
95% confidence intervals

```
dfNew = data.frame(year=15)
predict(lmRunVol, newdata = dfNew, interval = 'confidence')
```

```
##          fit          lwr          upr
## 1 10000.37 9261.564 10739.18
```

we can see below that the IQR of running times throughout the years has remained relatively constant. That means that

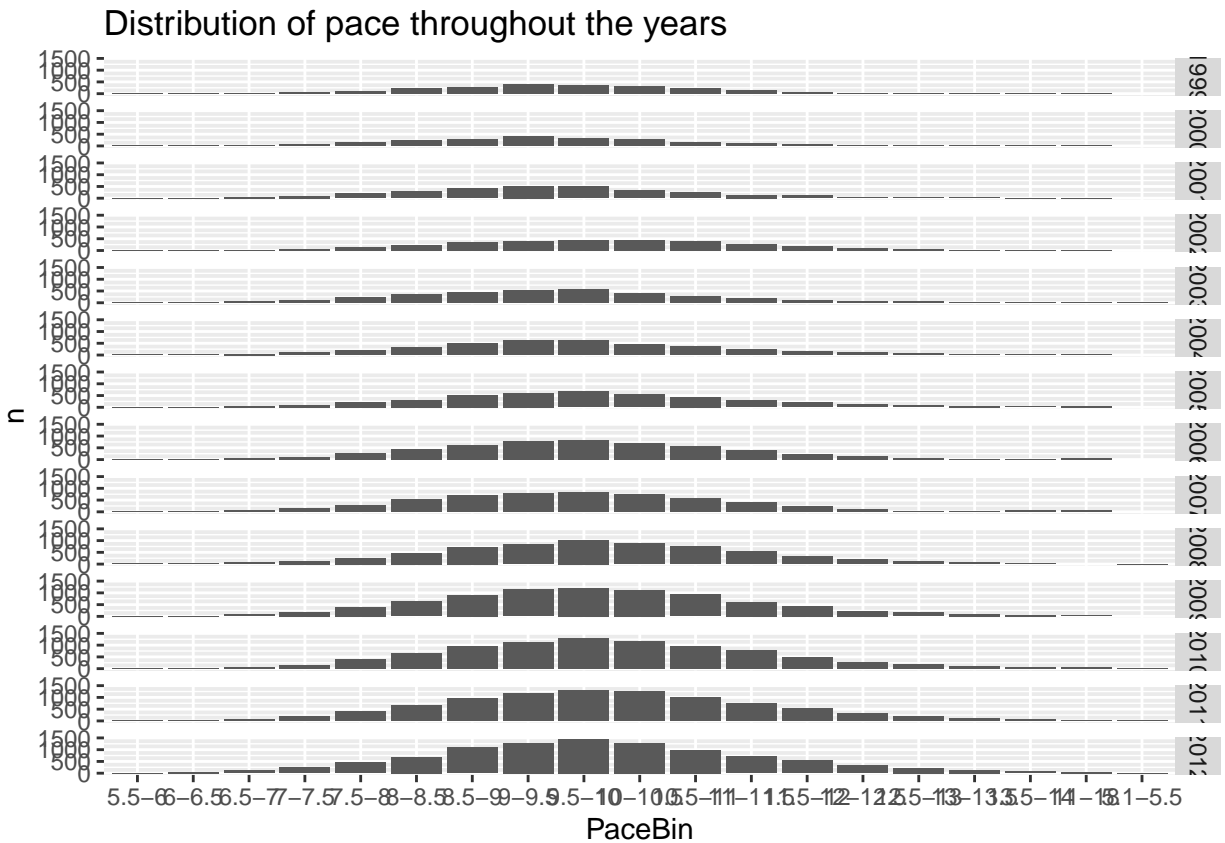
```
f = ggplot(womens_table_T)
f + geom_boxplot(mapping = aes(x=year, y=RawTime_M)) + theme_classic() + labs(title="Boxplot of Running
```



Let's look at the volume of runners by pace bin. We can see that the majority of runners are between 8-11 minute pace. This means that a large swell of runners will be running through the course together and race support must have enough volunteers to support the increase of runners as they make their way through the course.

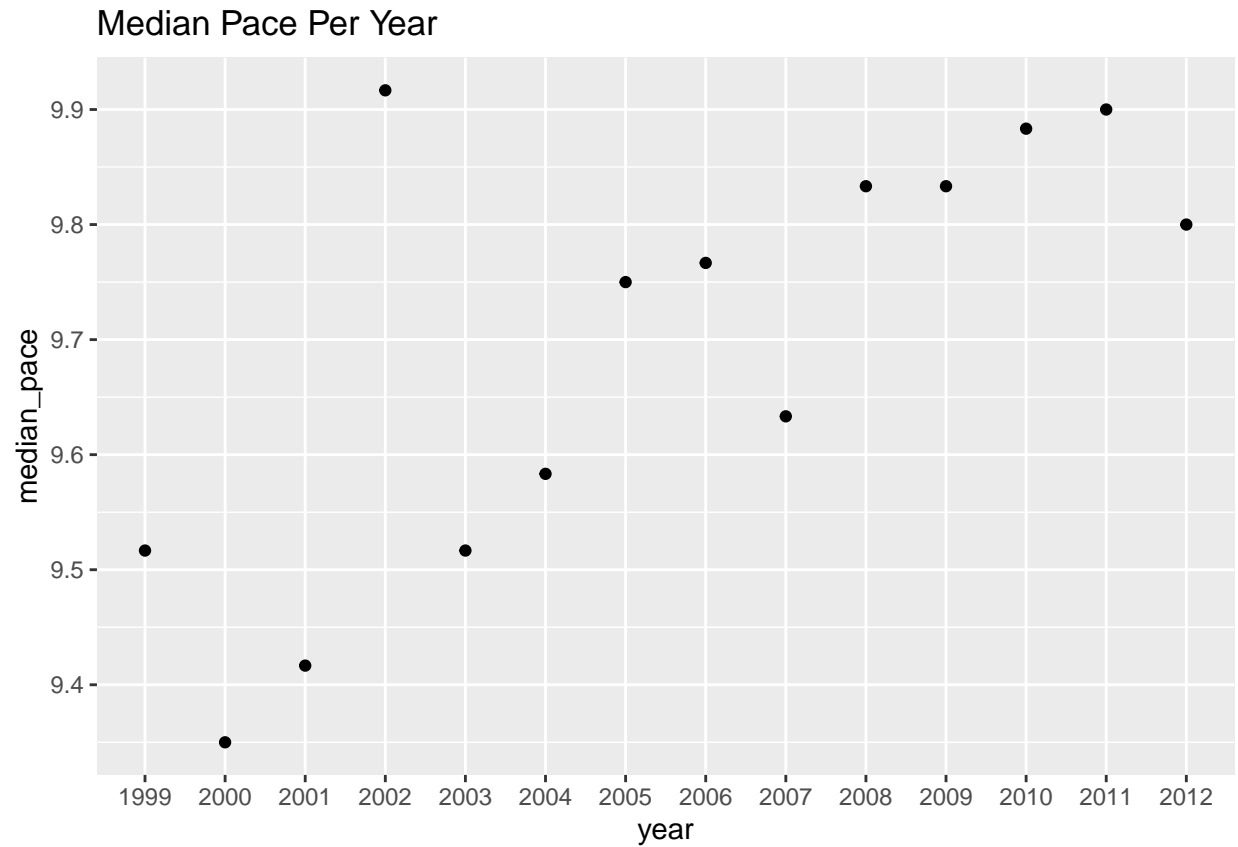
```
womens_table_T$PaceBin <- as.character(womens_table_T$PaceBin)
womens_table_T$PaceBin <- factor(womens_table_T$PaceBin, levels=unique(womens_table_T$PaceBin))

ggplot(womens_table_T %>% group_by(year, PaceBin) %>% tally()) +
  geom_col(mapping = aes(x=PaceBin, y=n)) +
  facet_grid(vars(year)) +
  labs(title="Distribution of pace throughout the years")
```




Let's take a look at the median pace per year.

```
womens_table_T %>% group_by(year) %>% dplyr::summarise(median_pace = median(RawPace_M)) %>% ggplot(aes(
## 'summarise()' ungrouping output (override with '.groups' argument)
```

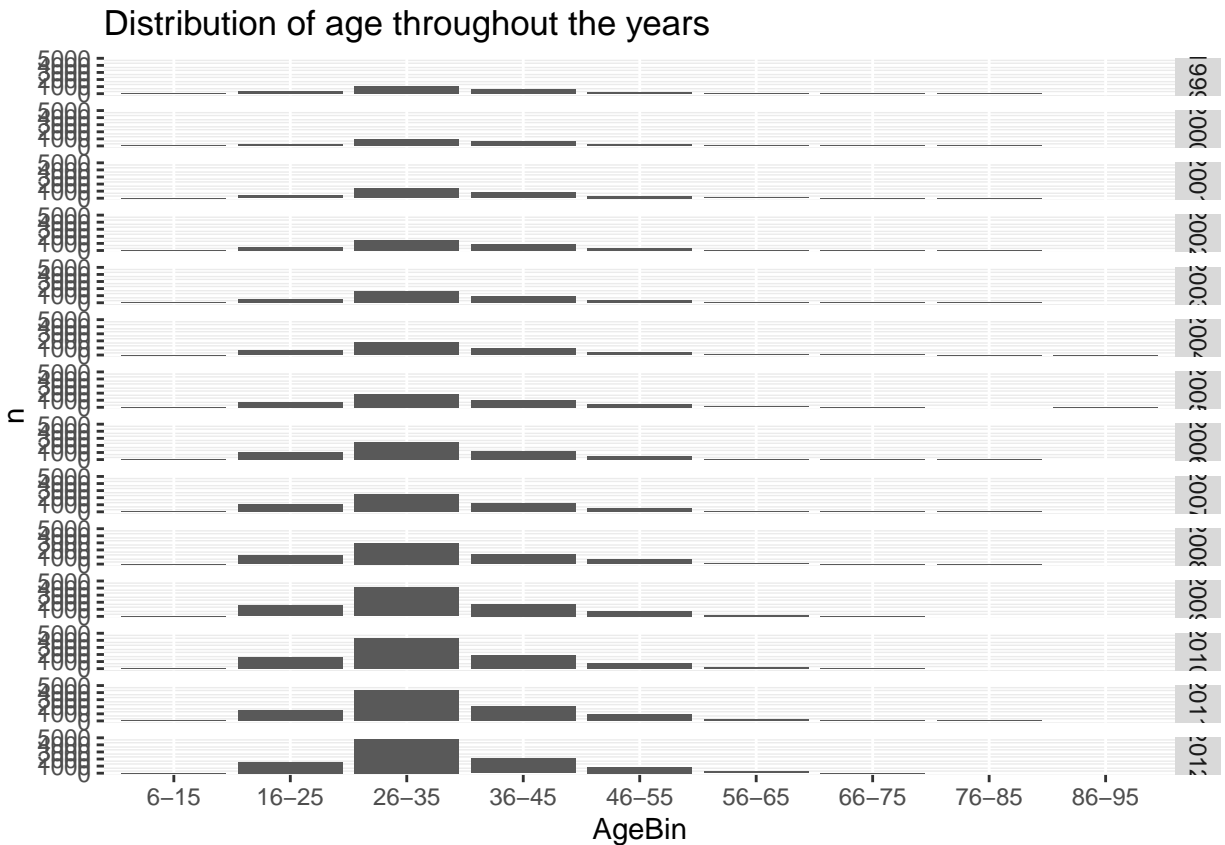
If we look at the max times throughout the years, there isn't a consistent pattern observable, we'll revisit this after looking more at the age bins further

```
womens_table_T %>% group_by(year) %>% dplyr::summarise(max_time = max( RawTime_M )) %>% ggplot(aes(x=year, y=max_time))  
  
## 'summarise()' ungrouping output (override with '.groups' argument)
```



we see below that there has been a steady increase in 26-35 year old runners over the years

```
ggplot(womens_table_T %>% group_by(year, AgeBin) %>% tally()) +  
  geom_col(mapping = aes(x=AgeBin, y=n)) +  
  facet_grid(vars(year)) +  
  labs(title="Distribution of age throughout the years")
```



Let's take a closer look at the volume of runners in the 26-35 age category

```
dfRunAgeBinVol = womens_table_T %>% group_by(year, AgeBin) %>% dplyr::summarise(AgeBinVolume = n())
```

```
## 'summarise()' regrouping output by 'year' (override with '.groups' argument)
```

```
head(dfRunAgeBinVol)
```

```
## # A tibble: 6 x 3
## # Groups:   year [1]
##   year AgeBin AgeBinVolume
##   <fct> <fct>         <int>
## 1 1999 6-15             5
## 2 1999 16-25          286
## 3 1999 26-35         1086
## 4 1999 36-45          649
## 5 1999 46-55          273
## 6 1999 56-65          43
```

```
ggplot(dfRunAgeBinVol[dfRunAgeBinVol$AgeBin == '26-35',]) +
  geom_point(mapping = aes(x=year, y=AgeBinVolume)) +
  #facet_wrap(vars(year)) +
  labs(title="Increase in volume of 26-35 age group throughout the years")
```




```
dfRunAgeBinVol$year = as.numeric(as.character(dfRunAgeBinVol$year))
dfRunAgeBinVol[dfRunAgeBinVol$AgeBin == '26-35',]
```

```
## # A tibble: 14 x 3
## # Groups:   year [14]
##   year AgeBin AgeBinVolume
##   <dbl> <fct>         <int>
## 1  1999 26-35           1086
## 2  2000 26-35            914
## 3  2001 26-35          1368
## 4  2002 26-35          1497
## 5  2003 26-35          1579
## 6  2004 26-35          1793
## 7  2005 26-35          1882
## 8  2006 26-35          2464
## 9  2007 26-35          2523
## 10 2008 26-35          2987
## 11 2009 26-35          4062
## 12 2010 26-35          4316
## 13 2011 26-35          4340
## 14 2012 26-35          4837
```

Let's trend the volume of 26-35 year old runners using LOESS

```
df26 = dfRunAgeBinVol[dfRunAgeBinVol$AgeBin == '26-35',]
loessAge26 = loess(AgeBinVolume ~year, data=df26)
preds = predict(loessAge26)
ggplot(df26, aes(x=year, y=AgeBinVolume)) +
  geom_point() +
  geom_line(df26, mapping = aes(x=year, y=preds, col = "red")) +
  ggtitle("Prediction of Runner Volume in 26-35 Age Category")
```



Let's now try to predict the volume of 26-35 year olds using a piece wise linear regression model

```
pw.model = piecewise.linear(df26$year, df26$AgeBinVolume, middle=1, CI=TRUE, sig.level = 0.05)
pw.model
```

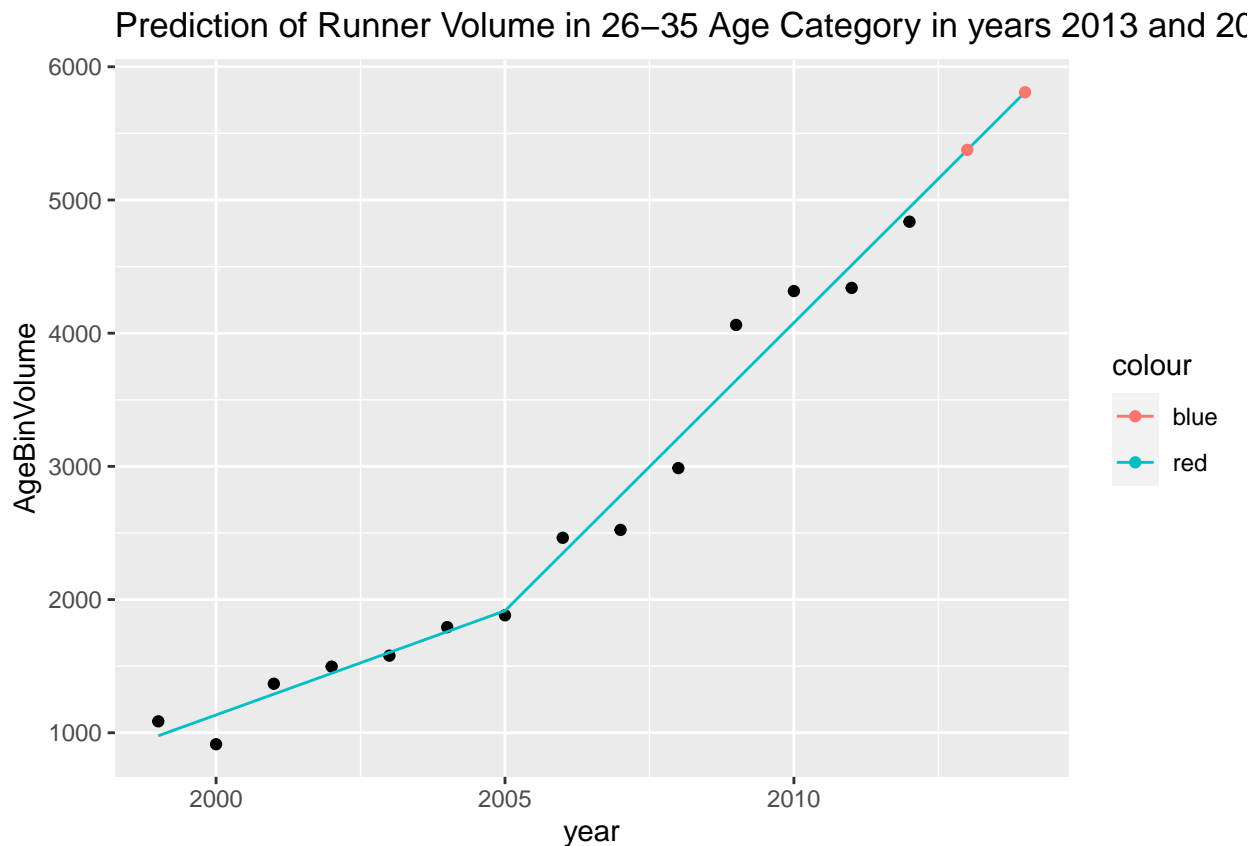
```
## [1] "Threshold alpha: 2005.0000564959"
## [1] ""
## [1] "Model coefficients: Beta[0], Beta[1], Beta[2]"
## (Intercept)          x          w
## -311648.2551    156.3910    276.1541
##      Change.Point Initial.Slope Slope.Change Second.Slope
## 2.5%      2003.007      81.99976    185.6956    369.3361
## 97.5%      2007.500     237.22169    570.7324    701.8154
```

Now that we have a piece wise linear model, let's predict the volume of 26-35 year olds in years 2013 and 2014

```

preds = predict(pw.model,c(1999:2014))
predYears = c(1999:2014)
dfPreds = data.frame(year = predYears, predAgeBinVolume = predict(pw.model,predYears))
ggplot(df26, aes(x=year, y=AgeBinVolume)) +
  geom_point() +
  geom_line(dfPreds, mapping = aes(x=year, y=predAgeBinVolume, col = "red")) +
  geom_point(dfPreds[dfPreds$year %in% c(2013:2014), ], mapping = aes(x=year, y=predAgeBinVolume, col = "blue")) +
  ggtitle("Prediction of Runner Volume in 26-35 Age Category in years 2013 and 2014")

```



The actual predicted values for years 2013 and 2014. This is what we want to share with the race management team

```
predict(pw.model,c(2013:2014))
```

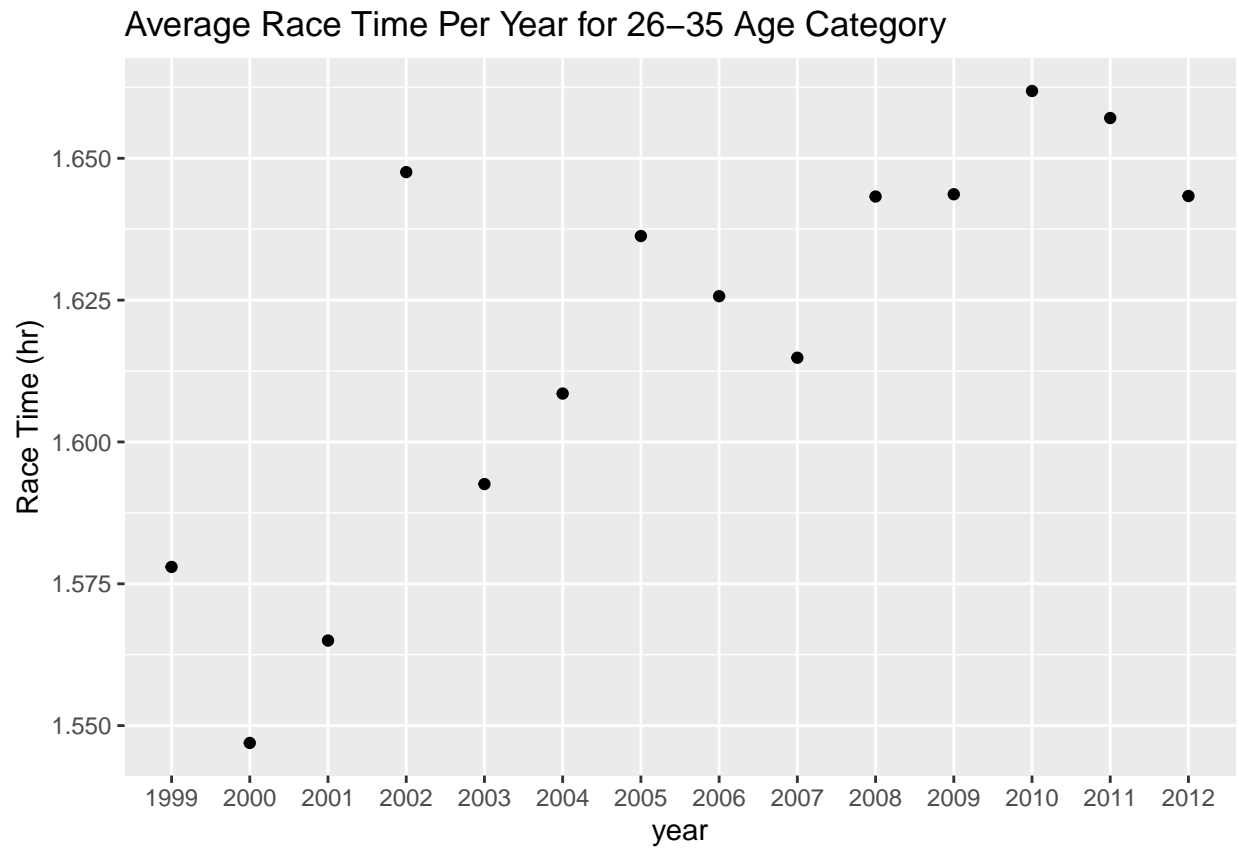
```
## [1] 5376.135 5808.681
```

Let's look at average times of age 25-36 year olds

```
dfAvgTimes26 = womens_table_T %>% group_by(year, AgeBin) %>% dplyr::summarise(avgTime = mean(RawTime_M))
```

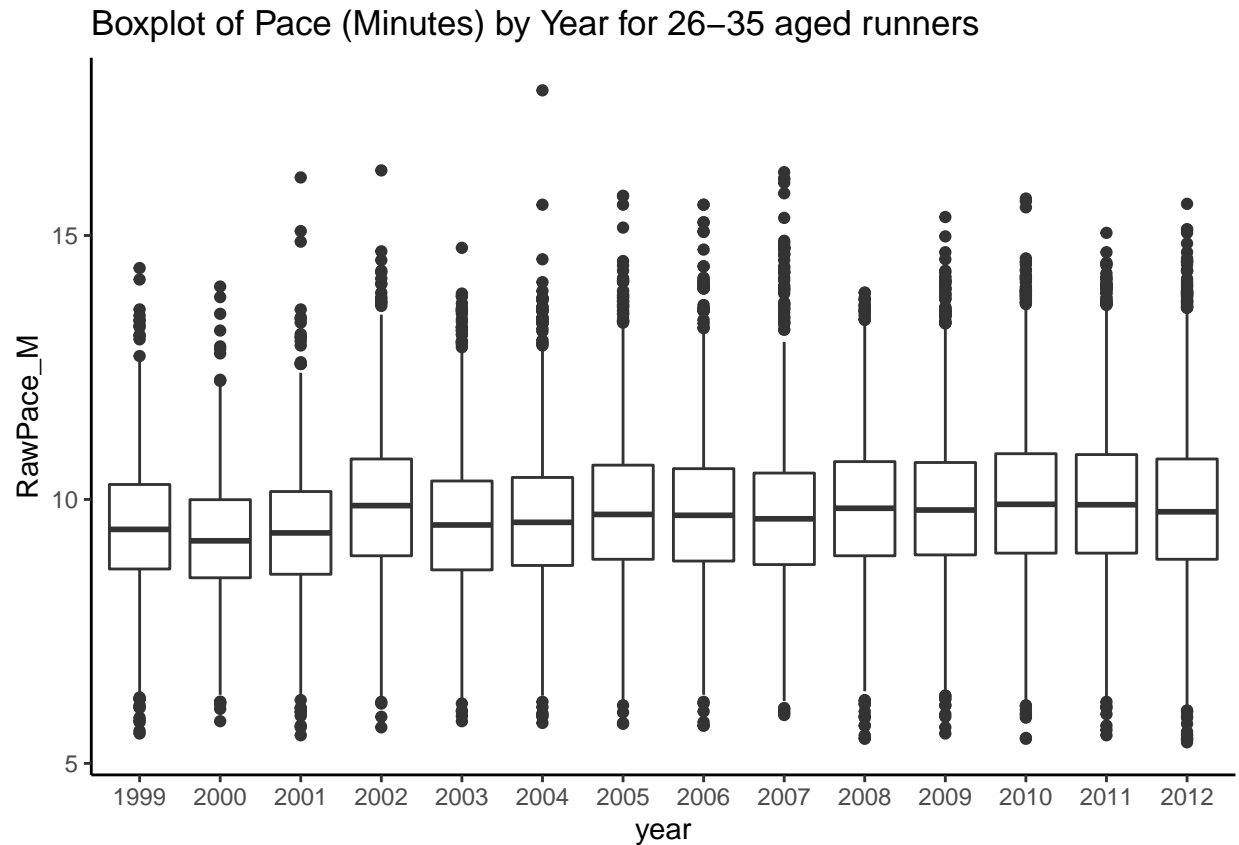
```
## 'summarise()' regrouping output by 'year' (override with '.groups' argument)
```

```
dfAvgTimes26 = dfAvgTimes26[dfAvgTimes26$AgeBin == '26-35',]
dfAvgTimes26 %>% ggplot(aes(x=year, y=avgTime/60)) + geom_point() + ggtitle("Average Race Time Per Year
```



let's get a look at the pace of 26-35 year olds. over the years, they've held a pretty consistent pace. 50% are in the 9-10 minute pace group.

```
f = ggplot(womens_table_T %>% filter(womens_table_T$AgeBin == "26-35"))
f + geom_boxplot(mapping = aes(x=year, y=RawPace_M)) + theme_classic() + labs(title="Boxplot of Pace (M
```

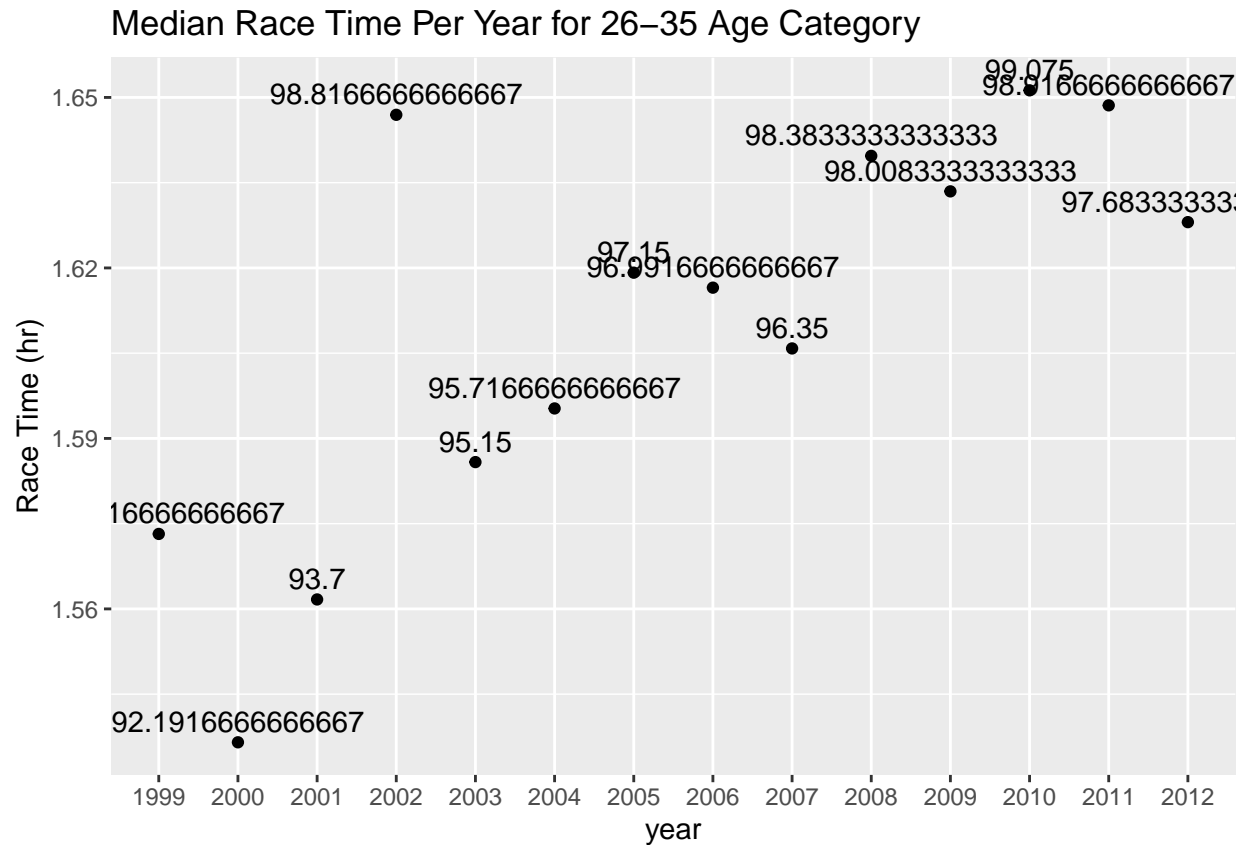


Let's look at median times of age 25-36 year olds

```
dfMedTimes26 = womens_table_T %>% group_by(year, AgeBin) %>% dplyr::summarise(medTime = median(RawTime_
```

```
## 'summarise()' regrouping output by 'year' (override with '.groups' argument)
```

```
dfMedTimes26 = dfMedTimes26[dfMedTimes26$AgeBin == '26-35',]
dfMedTimes26 %>% ggplot(aes(x=year, y=medTime/60)) + geom_point() + ggtitle("Median Race Time Per Year :")
geom_text(aes(label = medTime), vjust = -0.5)
```

Let's look at the proportion the 26-35 year olds make up of the total runners over the years. As of 2012, this age group makes up about 50% of all runners.

```
dfAgg = womens_table_T %>% filter(womens_table_T$AgeBin == "26-35") %>% group_by(year, AgeBin) %>% dplyr::summarise(
  TotalRunners = n())

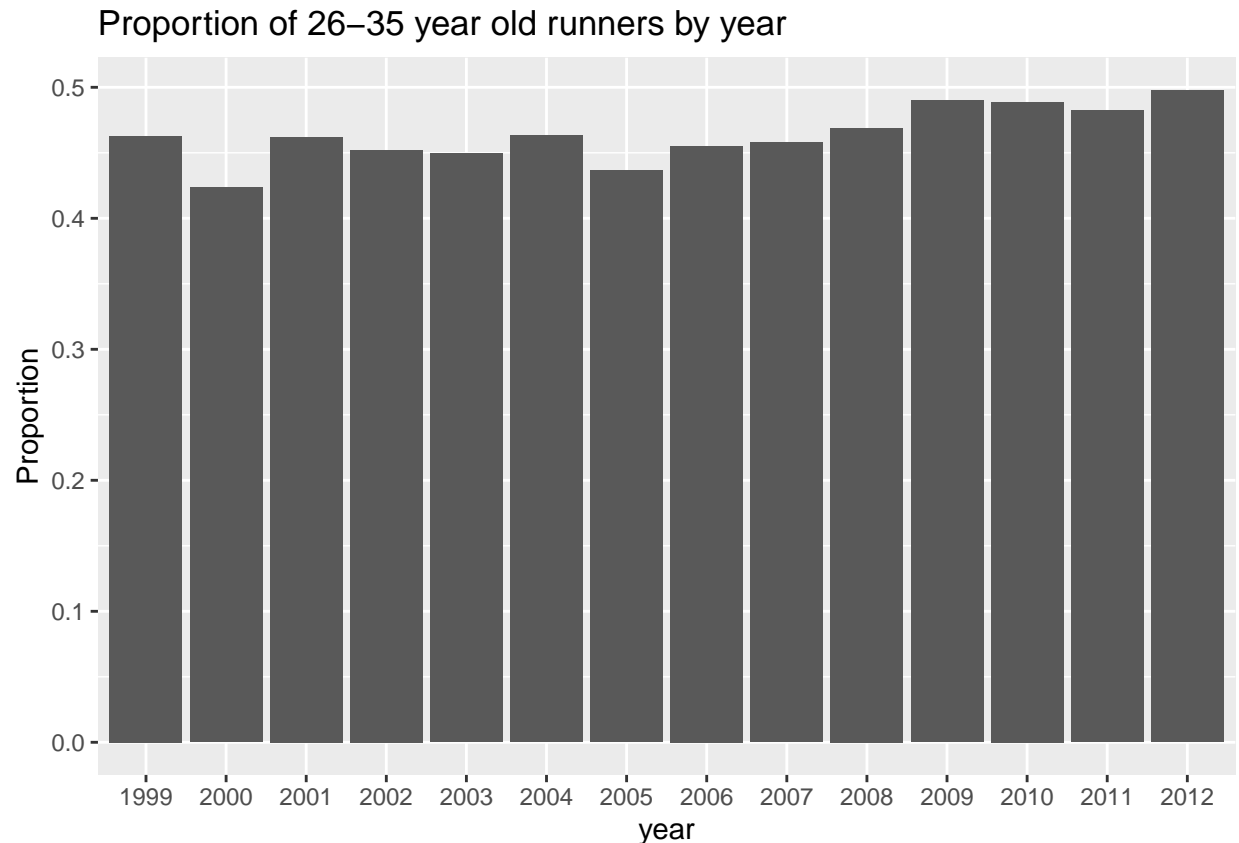
## 'summarise()' regrouping output by 'year' (override with '.groups' argument)

dfAgg$TotalRunners = (womens_table_T %>% group_by(year) %>% dplyr::summarise(TotalRunners = n()))$TotalRunners

## 'summarise()' ungrouping output (override with '.groups' argument)

dfAgg$Proportion = dfAgg$AgeBinVolume / dfAgg$TotalRunners

ggplot(dfAgg, aes(x=year, y=Proportion)) +
  geom_col() +
  labs(title="Proportion of 26-35 year old runners by year")
```




So to summarize, it's expected that 50% of the runners will be in the 26-35 year old age group in 2013. The median race time will be 97-98 minutes or approximately 9.7 minute pace.

Sabrina add on#####

When we consider how things have changed, we were asked to look at the average racer time. For that, we are going to first consider a LOESS prediction. Because this is a nonparametric method, and will weight itself relative to time, we feel like this should give indication of whether or not the average runner is getting faster or slower.

```
#Predict LOESS on time
womens_table_T$RawTime_M = as.numeric(womens_table_T$RawTime_M) #need data to be numeric for LOESS
womens_table_T$year = as.numeric(womens_table_T$year) #need data to be numeric for LOESS
plot(RawTime_M~year, ylim = c(94,100), data=womens_table_T, main="Year V Time")
out <- loess(RawTime_M~year, data=womens_table_T)
curve(predict(out, newdata=data.frame(year = x)), add=TRUE)
```



As shown above average time actually increased across the 14 year window, but did show a downturn in 2011 and 2012. This indicates that the finishers may in fact be speeding back up again.

Our next objective is to determine if anywhere along the way the average time has sectioned off in a way that indicates multiple trends. For that, we will run a changepoint evaluation on the average.

```
#Run a piecewise fit next
plotdata = womens_table_T %>%
  group_by(year) %>%
  summarise(average = mean(RawTime_M), .groups = "keep")

plotdata$average = as.integer(plotdata$average)
plotdata$average = as.ts(plotdata$average)

library(changepoint)
```

```
## Loading required package: zoo
```

```
##
```

```
## Attaching package: 'zoo'
```

```
## The following objects are masked from 'package:base':
```

```
##
```

```
## as.Date, as.Date.numeric
```


```
## Successfully loaded changepoint package version 2.2.2
```

```
## NOTE: Predefined penalty values changed in version 2.2. Previous penalty values with a postfix 1 i
```

```
dis.bs <- cpt.meanvar(plotdata$average, test.stat = "Poisson", method = "BinSeg")
pts.ts(dis.bs)
```

```
## numeric(0)
```

```
plot(dis.bs, cpt.width = 3)
```



You can see that this produced a single average. This could be appropriately stating that nothing has statistically changed. More likely, we do not have enough data points to find clean break points.

Conclusion / Recommendations

As mentioned in previous sections, the average age of the female runner has decreased in the last 14 years. The average race time of participants has slowed, and the distribution has formed a wider, more even presentation. From this, we conclude that participation in this race has become more popular, and likely as much a social event as a competitive one. As you look to advertise in the future, you should consider expansion of advertising into social media platforms, where you are likely to connect the most effectively with the demographic described above.

As with any expansion, planning must be taken into consideration for an enjoyable event. In the 14 years we evaluated, your race participation size in females quadrupled. We predict you will see over 10,000 female participants in the next year's event. You will need to consider the logistics in organizing to that scale. You should anticipate needing 10% more staff on race day. Volunteers to pass out water, hand out medals, and manage bag checks are one part of that increase. You'll also need additional police to help maintain crowd control. If you provide pace runners, you should consider that in some of your highest volume pace groups,

you may need an additional runner per assigned segment to ensure your participants are able to manage their positions.

Your site provides contact information for hotel accommodations booked at group rates. In the last four years, out of state participation has doubled (based on excluding DC, VA, and MD). Knowing that we expect participation to be higher next year, you should factor this into any hotel negotiations for the next event. You must also ensure that you have sufficient parking, or shuttle service in place if you unable to allow participants and fans to park directly adjacent to the starting line.

```
# looking at counts with division splits too
plotdata = womens_table_T %>% filter(HomeState != "VA") %>% filter(HomeState != "MD") %>% filter(HomeState != "DC")
  group_by(year, HomeState) %>%
  summarise(count = n())
```

```
## 'summarise()' regrouping output by 'year' (override with '.groups' argument)
```

```
p = plotdata %>%
  ggplot(aes(x = year, y = count, fill = HomeState)) +
  geom_bar(stat = "identity", position = "stack")
ggplotly(p)
```

Finally, we recommend that you have contingencies for the unexpected. A myriad of situations could cause this event to be delayed, postponed and even cancelled. Coordinating 10,000 people converging on Washington DC is no small feat, and that is only your female participation. Should you encounter a situation that makes it unsafe to gather, you need to have options available to participants. If you entertain hosting a virtual run in lieu of gathering in person, you must consider the supply chain impacts of needing to distribute race swag and medals. Ordering well in advance, and in bulk since you know your expected participation, will allow you the time needed to handle everyone's participation memorabilia one extra time.